

# SQL For Testers

Introduction

Importance of Databases and SQL

# Database Components

Tables, Schemas, Columns, and Keys

Constraints and Views

Stored Procedures and Triggers

# Structuring the Database

Introducing Database Normalization

Advancing Database Normalization

SQL Data Types

# Data Retrieval

Query Execution Process

Guidelines for Writing Query

Basic Query Syntax

Learn Data Retrieval Through Practice

# Data Definition

Data Definition Language (DDL)

Learn DDL Through Practice

Running the DDL Script

# Data Manipulation

Data Manipulation Language (DML)

Learn DML Through Practice

Running the DML Script

# Transforming Data and Functions

## Aggregates and Joins

Learn Aggregates and Joins Through Practice

Data Presentation

Learn Data Presentation Through Practice

Data Transformation

Learn Data Transformation Through Practice



# Database Administration

Backup, Recovery, and Database Restore

Learn Backup, Recovery, and Restore Through Practice

Login, Users, and Roles

Learn Server Logins, Users, and Roles Through Practice

# Understand Normalization Rules

Testing database normalization involves verifying that your database schema adheres to the rules of normalization.

While there isn't a direct automated tool to check normalization, you can perform manual checks and write scripts to validate your schema. Let's explore some practical steps:

### First Normal Form (1NF):

Ensure each column contains atomic (indivisible) values.

Write a script to check if any columns have multiple values (e.g., comma-separated lists).

### Second Normal Form (2NF):

Verify that non-key attributes depend fully on the entire primary key.

Write a script to identify partial dependencies.

### Third Normal Form (3NF):

Check for transitive dependencies.

Write a script to find columns that depend on other non-key attributes.

Run Validation Queries:

Execute your scripts against your database.

Identify any violations of normalization rules.

Automate Checks (Optional):


# The normal forms

*V. dangerous*



*Safety*

*V. safe*



*First normal form*

*Second normal form*

*Third normal form*

*Fourth normal form*

*Fifth normal form*

Who were the members  
of the Beatles?

John, Paul, George, and Ringo

Paul, John, Ringo, and George

(equivalent)

Members of the Beatles  
from Tallest  
to Shortest

Paul

John

George

Ringo

Members of the Beatles  
from Tallest  
to Shortest

Paul

John

George

✓ Ringo

Not normalized

Using row order to convey  
information violates 1NF.

## Beatle\_Height

Beatle	Height_In_Cm
George	178
John	179
Ringo	170
Paul	180



## Beatle\_Height

Beatle	Height_In_Cm
George	178
John	179
Ringo	Somewhere between 168 and 171
Paul	180

## Beatle\_Height

Beatle	Height_In_Cm (integer)
George	178
John	179
Ringo	Somewhere between 168 and 171
Paul	180

Mixing data types within the same column violates 1NF

(and the DB platform won't let you do it anyway)

We should make this  
the primary key...



Beatle\_Height

Beatle	Height_In_Cm
George	178
John	179
Ringo	170
Paul	180

```
ALTER TABLE Beatle_Height  
ADD PRIMARY KEY (Beatle);
```

Making "Beatle" the primary key prevents this:

Beatle	Height_In_Cm
George	178
John	

Player_ID	Inventory
jdog21	2 amulets, 4 rings
gilal9	18 copper coins
trev73	3 shields, 5 arrows, 30 copper coins, 7 rings

## Player\_Inventory

Player_ID	Quantity_1	Item_Type_1	Quantity_2	Item_Type_2	Quantity_3	Item_Type_3	Quantity_4	Item_Type_4
jdog21	2	amulets	4	rings				
gilal9	18	copper coins						
trev73	3	shields	5	arrows	30	copper coins	7	rings



## Player\_Inventory

Player_ID	Item_Type	Item_Quantity
jdog21	amulets	2
jdog21	rings	4
gilal9	copper coins	18
trev73	shields	3
trev73	arrows	5
trev73	copper coins	30
trev73	rings	7

# First Normal Form Rules:

- 1) Using row order to convey information is not permitted
- 2) Mixing data types within the same column is not permitted
- 3) Having a table without a primary key is not permitted
- 4) Repeating groups are not permitted

## 2nd Normal Form Rules



# 3rd Normal Form Rules

## 4th Normal Form Rules

## 5th Normal Form Rules

# Querying data

- `SELECT FROM` – show you how to use a simple `SELECT FROM` statement to query the data from a single table.
- `SELECT` – learn how to use the `SELECT` statement without referencing a table.

# Sorting data

- ORDER BY

# Filtering data

- WHERE – learn how to use the WHERE clause to filter rows based on specified conditions.
- SELECT DISTINCT – show you how to use the DISTINCT operator in the SELECT statement to eliminate duplicate rows in a result set.
- AND – introduce you to the AND operator to combine Boolean expressions to form a complex condition for filtering data.
- OR– introduce you to the OR operator and show you how to combine the OR operator with the AND operator to filter data.
- IN – show you how to use the IN operator in the WHERE clause to determine if a value matches any value in a set.
- NOT IN – negate the IN operator using the NOT operator to check if a value doesn't match any value in a set.
- BETWEEN – show you how to query data based on a range using the BETWEEN operator.
- LIKE – query database on pattern matching using wildcards such as % and \_.
- LIMIT – use LIMIT to limit the number of rows returned by the SELECT statement
- IS NULL – test whether a value is NULL or not by using the IS NULL operator.

# Joining tables

- Table & Column Aliases – introduce you to table and column aliases.
- Joins – give you an overview of joins supported in MySQL including inner join, left join, and right join.
- INNER JOIN – query rows from a table that has matching rows in another table.
- LEFT JOIN – return all rows from the left table and matching rows from the right table or null if no matching rows are found in the right table.
- RIGHT JOIN – return all rows from the right table and matching rows from the left table or null if no matching rows are found in the left table.
- Self-join – join a table to itself using a table alias and connect rows within the same table using inner join and left join.
- CROSS JOIN – make a Cartesian product of rows from multiple tables.

# WHY JOINS?

- Typically in a relational database, data is organized into various tables made of attributes (columns) and records (rows).
- In each table there exist a column that is the **primary key** which is a column where each entry uniquely represents a single row in that table. This is usually the ID (short for identifier) column.
- A column in a table that establishes an association with another table's primary key via shared values is called a **foreign key**. Foreign keys are also typically titled IDs but prepended with the name of the referenced table.
- This concept is applied when combining two or more tables together using a JOIN

	lt_id	lt-data1
▶	1	LEFT-DATA1
	2	LEFT-DATA2
	3	LEFT-DATA3

	rt_id	rt_fk	rt_data1
▶	1	1	RIGHT-DATA1
	2	1	RIGHT-DATA2
	3	4	RIGHT-DATA3
	4	2	RIGHT-DATA4




User Table – Table 1

ID (Primary Key)	Name	Address
1	Sally Select	123 Join Dr
2	Frank From	25 Where St

Event Table – Table 2

User_ID (Foreign Key)	ID (Primary Key)	Action
1	A	LOGIN
3	B	VIEW PAGE
4	C	LOGIN

# Outer Join

Table 1 

1		
2		

Table 2 

1		
3		
4		

Outer Join 

1				
2				
3				
4				

# Inner Join

```
SELECT Orders.OrderID, Customers.CustomerName, Orders.OrderDate
```


```
FROM
```

```
Orders INNER JOIN Customers ON Orders.CustomerID=Customers.CustomerID;
```

**INNER JOIN: This returns records that have matching values in both tables**

Table 1 

1		
2		

Table 2 

1		
3		
4		


Inner Join 

1				

# Left Join

Table 1 

1		
2		


Table 2 

1		
3		
4		


Left Join 

1				
2				

# Cross Join

Table 1 

1		
2		

Table 2 

1		
3		
4		

Cross Join 

1			1		
1			3		
1			4		
2			1		
2			3		
2			4		

# Joins Summary

## Combining Data Tables – SQL Joins Explained

A JOIN clause in SQL is used to combine rows from two or more tables, based on a **related column** between them.

Table 1 

1		
2		

Table 2 

1		
3		
4		

Outer Join 

1				
2				
3				
4				

Inner Join 

1				
---	--	--	--	--

Left Join 

1				
2				

Union  + 

1		
2		
1		
3		
4		

Cross Join 

1			1		
1			3		
1			4		
2			1		
2			3		
2			4		

```
SELECT * FROM friends WHERE name = 'Zack';
```

Indexing makes columns faster to query by creating pointers to where data is stored within a database. Imagine you want to find a piece of information that is within a large database. To get this information out of the database the computer will look through every row until it finds it. If the data you are looking for is towards the very end, this query would take a long time to run.

This took 3 comparisons to find the right answer instead of 8 in the unindexed data.

Indexes allow us to create sorted lists without having to create all new sorted tables, which would take up a lot of storage space.

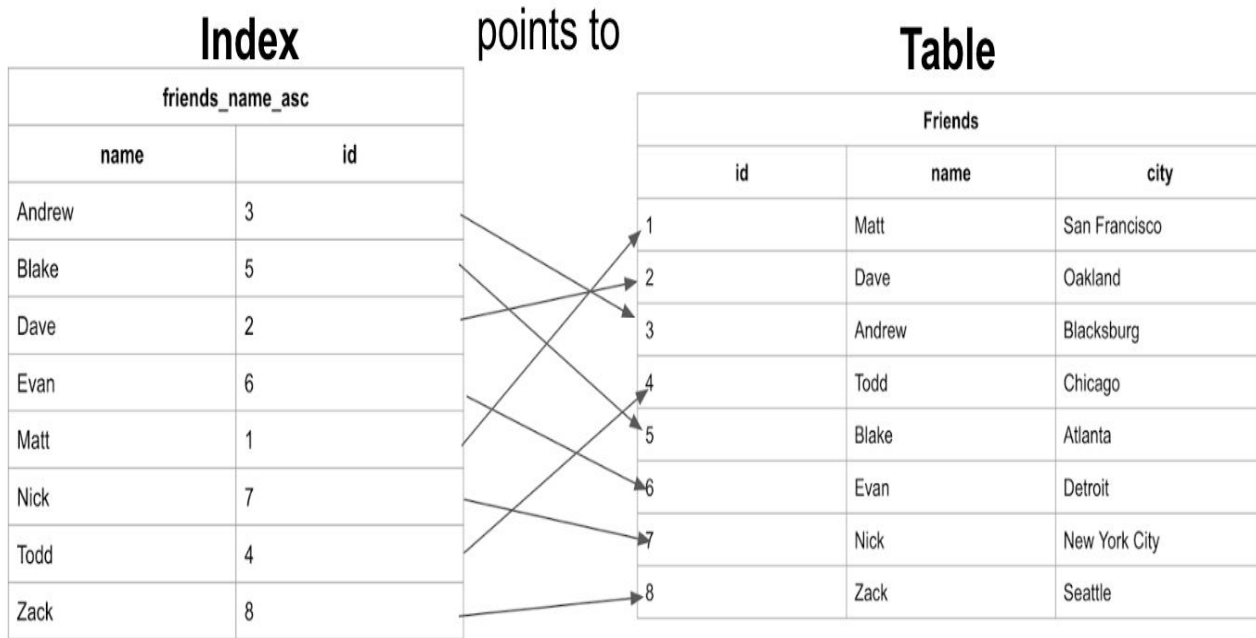
friends		
id	name	city
1	Matt	San Francisco
2	Dave	Oakland
3	Andrew	Blacksburg
4	Todd	Chicago
5	Blake	Atlanta
6	Evan	Detroit
7	Nick	New York City
8	Zack	Seattle

```
SELECT * FROM friends WHERE name = 'Zack';
```

friends_name_asc	
Name	Index
Andrew	3
Blake	5
Dave	2
Evan	6
Matt	1
Nick	7
Todd	4
Zack	8

# Whats is an Index

An index is a structure that holds the field the index is sorting and a pointer from each record to their corresponding record in the original table where the data is actually stored.





# Types of Indexes

There are two types of databases indexes:

## Clustered

- Clustered indexes are the unique index per table that uses the primary key to organize the data that is within the table.
- Clustered indexes do not have to be explicitly declared.
- Created when the table is created.
- Use the primary key sorted in ascending order.

## Non-clustered

- Non-clustered indexes are sorted references for a specific field, from the main table, that hold pointers back to the original entries of the table.

```
CREATE [UNIQUE | FULLTEXT | SPATIAL] INDEX index_name
    [index_type]
    ON tbl_name (key_part,...)
    [index_option]
    [algorithm_option | lock_option] ...
```

*key\_part*: {*col\_name* [(*length*)] | (*expr*)} [ASC | DESC]

```
index_option: {
    KEY_BLOCK_SIZE [=] value
| index_type
| WITH PARSER parser_name
| COMMENT 'string'
| {VISIBLE | INVISIBLE}
| ENGINE_ATTRIBUTE [=] 'string'
| SECONDARY_ENGINE_ATTRIBUTE [=] 'string'
}
```

```
index_type:
    USING {BTREE | HASH}
```

```
algorithm_option:
    ALGORITHM [=] {DEFAULT | INPLACE | COPY}
```

```
lock_option:
    LOCK [=] {DEFAULT | NONE | SHARED | EXCLUSIVE}
```

# Summary

---

- Indexing can vastly reduce the time of queries
- Every table with a primary key has one clustered index
- Every table can have many non-clustered indexes to aid in querying
- Non-clustered indexes hold pointers back to the main table
- Not every database will benefit from indexing
- Not every index will increase the query speed for the database